



國立陽明交通大學

NATIONAL YANG MING CHIAO TUNG UNIVERSITY

Deep Learning  
Lab6: DQN and DDPG

Name: 許子駿

Student ID: 311551166

Institute of Computer Science and Engineering

2023 年 5 月 27 日

# Table of Contents

|          |   |           |
|----------|---|-----------|
| <b>1</b> | <b>Introduction</b>   | <b>2</b>  |
| <b>2</b> | <b>Experiment setups</b>  | <b>3</b>  |
| 2.1      | DQN . . . . .   | 3         |
| 2.1.1    | Network . . . . .   | 3         |
| 2.1.2    | Action selection . . . . .  | 4         |
| 2.1.3    | Updating behavior network . . . . .   | 4         |
| 2.1.4    | Updating target network . . . . .   | 5         |
| 2.2      | DDPG . . . . .  | 5         |
| 2.2.1    | Actor network . . . . .   | 5         |
| 2.2.2    | Critic network . . . . .  | 6         |
| 2.2.3    | Action selection . . . . .  | 6         |
| 2.2.4    | Updating behavior network . . . . .   | 7         |
| 2.2.5    | Updating target network . . . . .   | 8         |
| <b>3</b> | <b>Experimental results</b>   | <b>9</b>  |
| 3.1      | DQN for <i>LunarLander-v2</i> . . . . .   | 9         |
| 3.2      | DDPG for <i>LunarLanderContinuous-v2</i> . . . . .  | 10        |
| 3.3      | DQN for <i>BreakoutNoFrameskip-v4</i> . . . . .   | 10        |
| <b>4</b> | <b>Discussion</b>   | <b>12</b> |
| 4.1      | Explain effects of the discount factor . . . . .  | 12        |
| 4.2      | Explain benefits of epsilon-greedy in comparison to greedy action<br>selection . . . . .                                    | 12        |
| 4.3      | Explain the necessity of the target network . . . . .   | 12        |
| 4.4      | Describe the tricks you used in Breakout and their effects, and how<br>they differ from those used in LunarLander . . . . . | 13        |

# Chapter 1

## Introduction

This task is to solve *LunarLander-v2* and *BreakoutNoFrameskip-v4* using **deep Q-network (DQN)**, and *LunarLanderContinuous-v2* using **deep deterministic policy gradient (DDPG)**.

# Chapter 2

## Experiment setups

Experiment setups are listed in this chapter including 2 parts, DQN and DDPG.

### 2.1 DQN

DQN details are listed in this section including 4 parts, network, action selection, updating behavior network, and updating target network.

#### 2.1.1 Network

Since *LunarLander-v2* includes 4 actions, i.e., 0 (No-op), 1 (Fire left engine), 2 (Fire main engine), 3 (Fire right engine), a network with output dimension 4 is built to predict Q value (See Listing 2.1).

```
1 class Net(nn.Module):
2     def __init__(self,
3         state_dim=8, action_dim=4, hidden_dim=(400, 300)) -> None:
4         super().__init__()
5
6         self.layers = nn.Sequential(
7             nn.Linear(in_features=state_dim,
8                 out_features=hidden_dim[0]),
9             nn.ReLU(inplace=True),
10            nn.Linear(in_features=hidden_dim[0],
11                out_features=hidden_dim[1]),
12            nn.ReLU(inplace=True),
13            nn.Linear(in_features=hidden_dim[1],
14                out_features=action_dim)
15        )
16
```

```

17     def forward(self, x: torch.FloatTensor) -> torch.FloatTensor:
18         return self.layers(x)

```

Listing 2.1: Python code of `Net` of DQN.

### 2.1.2 Action selection

During the play, action with maximum Q value with probability  $\epsilon$  will be selected, or random action is selected, otherwise. It's called  $\epsilon$ -greedy (See Listing 2.2).

```

1 def select_action(self, state, epsilon, action_space):
2     if random.random() > epsilon:
3         with torch.no_grad():
4             # Max element is in 2nd column (dim=1).
5             state = torch.from_numpy(state).view(1, -1).to(self.device)
6             return self._behavior_net(state).max(dim=1)[1].item()
7     else:
8         return action_space.sample()

```

Listing 2.2: Python code of `select_action` of DQN.

### 2.1.3 Updating behavior network

To update behavior network, sample transitions (state, action, reward, next action, done) from replay memory, then do TD-learning, and finally calculate MSE loss from Q value and Q target (See Listing 2.3 and Equation 2.1).

```

1 def _update_behavior_network(self, gamma):
2     state, action, reward, next_state, done = self._memory.sample(
3         self.batch_size, self.device)
4
5     q_value = self._behavior_net(state).gather(dim=1, index=action.long())
6     with torch.no_grad():
7         q_next = self._target_net(next_state).max(dim=1)[0].view(-1, 1)
8         q_target = reward + gamma * q_next * (1 - done)
9
10    criterion = nn.MSELoss()
11    loss = criterion(q_value, q_target)
12
13    self._optim.zero_grad()
14    loss.backward()
15    nn.utils.clip_grad_norm_(self._behavior_net.parameters(), 5)
16    self._optim.step()

```

Listing 2.3: Python code of `_update_behavior_network` of DQN.

$$y_j = \begin{cases} r_j & \text{if episode terminates at step } j + 1 \\ r_j + \gamma \max_{a'} \hat{Q}(\phi_{j+1}, a'; \theta) & \text{otherwise} \end{cases} \quad (2.1)$$

### 2.1.4 Updating target network

To update target network, copy weights of behavior network (See Listing 2.4).

```
1 def _update_target_network(self):
2     self._target_net.load_state_dict(self._behavior_net.state_dict())
```

Listing 2.4: Python code of `_update_target_network` of DQN.

## 2.2 DDPG

DDPG details are listed in this section including 5 parts, actor network, critic network, action selection, updating behavior network, and updating target network.

### 2.2.1 Actor network

Since *LunarLanderContinuous-v2* includes 2 engines, i.e., main engine with range  $[-1, 1]$  and left-right engine with range  $[-1, 1]$ , a network with output dimension 2 is built (See Listing 2.5).

```
1 class ActorNet(nn.Module):
2     def __init__(self,
3         state_dim=8, action_dim=2, hidden_dim=(400, 300)) -> None:
4         super().__init__()
5
6         self.layers = nn.Sequential(
7             nn.Linear(in_features=state_dim,
8                 out_features=hidden_dim[0]),
9             nn.ReLU(inplace=True),
10            nn.Linear(in_features=hidden_dim[0],
11                out_features=hidden_dim[1]),
12            nn.ReLU(inplace=True),
13            nn.Linear(in_features=hidden_dim[1],
14                out_features=action_dim),
15            nn.Tanh()
16        )
17
18     def forward(self, x: torch.FloatTensor) -> torch.FloatTensor:
```

```
19 return self.layers(x)
```

Listing 2.5: Python code of **ActorNet** of DDPG.

### 2.2.2 Critic network

Critic network is used to predict Q value. Since the output is scalar, output dimension 1. (See Listing 2.6).

```
1 class CriticNet(nn.Module):
2     def __init__(self,
3         state_dim=8, action_dim=2, hidden_dim=(400, 300)) -> None:
4         super().__init__()
5
6         h1, h2 = hidden_dim
7         self.critic_head = nn.Sequential(
8             nn.Linear(state_dim + action_dim, h1),
9             nn.ReLU(),
10        )
11        self.critic = nn.Sequential(
12            nn.Linear(h1, h2),
13            nn.ReLU(),
14            nn.Linear(h2, 1),
15        )
16
17    def forward(self,
18        x: torch.FloatTensor, action: torch.FloatTensor) -> torch.
19    FloatTensor:
20        x = self.critic_head(torch.cat([x, action], dim=1))
21        return self.critic(x)
```

Listing 2.6: Python code of **CriticNet** of DDPG.

### 2.2.3 Action selection

During the play, select action with actor network and add extra noise (See Listing 2.7).

```
1 def select_action(self, state, noise=True):
2     with torch.no_grad():
3         if noise:
4             re = self._actor_net(torch.from_numpy(state).view(1, -1).to(self.
5                 device)) + \
6                 torch.from_numpy(self._action_noise.sample()).view(1, -1).to(
7                     self.device)
```

```

6         else:
7             re = self._actor_net(torch.from_numpy(state).view(1, -1).to(self.
device))
8
9         return re.to('cpu').numpy().squeeze()

```

Listing 2.7: Python code of `select_action` of DDPG.

## 2.2.4 Updating behavior network

To update behavior network, sample transitions (state, action, reward, next action, done) from replay memory, and calculate MSE loss from Q value of behavior network and Q target of target network to update critic network. Also, to maximize Q value by updating actor network  $\mu$ , calculate negative Q value expectation  $\mathbb{E}[-Q(s, \mu(s))]$  (See Listing 2.8 and Equation 2.2).

```

1 def _update_behavior_network(self, gamma):
2     # Sample transitions batch.
3     state, action, reward, next_state, done = self._memory.sample(
4         self.batch_size, self.device)
5
6     # Update critic.
7     q_value = self._critic_net(state, action)
8     with torch.no_grad():
9         a_next = self._target_actor_net(next_state)
10        q_next = self._target_critic_net(next_state, a_next)
11        q_target = reward + gamma * q_next * (1 - done)
12
13    criterion = nn.MSELoss()
14    critic_loss = criterion(q_value, q_target)
15
16    self._actor_net.zero_grad()
17    self._critic_net.zero_grad()
18    critic_loss.backward()
19    self._critic_opt.step()
20
21    # Update actor.
22    action = self._actor_net(state)
23    actor_loss = -self._critic_net(state, action).mean()
24
25    self._actor_net.zero_grad()
26    self._critic_net.zero_grad()
27    actor_loss.backward()

```



```
28 self._actor_opt.step()
```

Listing 2.8: Python code of `__update_behavior_network` of DDPG.

$$y_i = r_i + \gamma Q'(s_{t+1}, \mu'(s_{t+1} | \theta^{\mu'}) | \theta^{Q'}) \quad (2.2)$$

### 2.2.5 Updating target network

To update target network, **soft copy** weights of behavior network (See Listing 2.9).

```
1 def _update_target_network(target_net, net, tau):  
2     for target, behavior in zip(target_net.parameters(), net.parameters()):  
3         target.data.copy_((1 - tau) * target.data + tau * behavior.data)
```

Listing 2.9: Python code of `__update_target_network` of DDPG.

# Chapter 3

## Experimental results

Results are listed in this chapter including 3 parts, DQN for *LunarLander-v2*, DDPG for *LunarLanderContinuous-v2*, and DQN for *BreakoutNoFrameskip-v4*.

### 3.1 DQN for *LunarLander-v2*

From Figure 3.2, we can see that the average testing reward is **255.89**.




Figure 3.1: Rewards of DQN for *LunarLander-v2* during training.

```
(dl) root@5b48b3907782:/workspace/lab6# xvfb-run -s "-screen 0 1400x900x24" python dqn.py --mode test --device 1 --resumeckpt models/230520_085345_dqn-lunar-lander-v2_batch_size128_episodes1200warmup10000_lr0.0005gamma0.99decay0.995min0.01_freq-behavior4freq-target100_eps-eval0.001/dqn.pth
Namespace(mode='test', devices=[1], batch_size=128, num_workers=8, capacity=10000, num_train_episodes=1200, num_steps_warmup=10000, lr=0.0005, gamma=0.99, epsilon_decay=0.995, epsilon_min=0.01, freq_update_behavior=4, freq_update_target=100, seed=20200519, resume_ckpt='models/230520_085345_dqn-lunar-lander-v2_batch_size128_episodes1200warmup10000_lr0.0005gamma0.99decay0.995min0.01_freq-behavior4freq-target100_eps-eval0.001/dqn.pth', epsilon_eval=0.001, dir_writer='./runs/', dir_model='./models/')

/opt/conda/envs/dl/lib/python3.9/site-packages/gym/logger.py:30: UserWarning: WARN: Box bound precision lowered by casting to float32
warnings.warn(colorize('%s: %s'%( 'WARN', msg % args), 'yellow'))
Currently using device cuda1...
Finish loading checkpoint from models/230520_085345_dqn-lunar-lander-v2_batch_size128_episodes1200warmup10000_lr0.0005gamma0.99decay0.995min0.01_freq-behavior4freq-target100_eps-eval0.001/dqn.pth...
Start testing...
Total reward: 255.13
Total reward: 276.14
Total reward: 280.01
Total reward: 277.24
Total reward: 293.5
Total reward: 33.45
Total reward: 270.18
Total reward: 280.21
Total reward: 299.92
Total reward: 293.04
Average Reward: 255.89
```

Figure 3.2: Rewards of DQN for *LunarLander-v2* during evaluation.

### 3.2 DDPG for *LunarLanderContinuous-v2*

From Figure 3.4, we can see that the average testing reward is **206.39**.




Figure 3.3: Rewards of DDPG for *LunarLanderContinuous-v2* during training.

```
(dl) root@5b48b39077821/workspace/lab6# xvfb-run -s "-screen 0 1400x900x24" python ddp.py --mode test --device 1 --resumeckpt models/238521_063715_ddpg-lunar-lander-continuous-v2_batch_size64_episodes1200warmup10000_lr0.001rc0.001gamma0.99tau0.005/ddpg.pth
Namespaces(namespace='test', devices=[1], batch_size=64, num_workers=8, capacity=500000, num_train_episodes=1200, num_steps_warmup=10000, lr_actor=0.001, lr_critic=0.001, gamma=0.99, tau=0.005, seed=28200519, resume_ckpt='models/238521_063715_ddpg-lunar-lander-continuous-v2_batch_size64_episodes1200warmup10000_lr0.001rc0.001gamma0.99tau0.005/ddpg.pth', dir_writer='./runs/', dir_model='./models/')

/opt/conda/envs/dl/lib/python3.9/site-packages/gym/logger.py:30: UserWarning: WARN: Box bound precision lowered by casting to float32
  warnings.warn(colorize('%s: %s' % ('WARN', msg % args), 'yellow'))
Currently using device cuda:1 ...
Finish loading checkpoint from models/238521_063715_ddpg-lunar-lander-continuous-v2_batch_size64_episodes1200warmup10000_lr0.001rc0.001gamma0.99tau0.005/ddpg.pth ...
Start testing ...
Total reward: 251.21
Total reward: 153.31
Total reward: 282.03
Total reward: 145.49
Total reward: 229.51
Total reward: 264.30
Total reward: 219.92
Total reward: 297.10
Total reward: -44.63
Total reward: 265.44
Average Reward: 206.39
```

Figure 3.4: Rewards of DDPG for *LunarLanderContinuous-v2* during evaluation.

### 3.3 DQN for *BreakoutNoFrameskip-v4*

From Figure 3.6, we can see that the average testing reward is **227.06**.

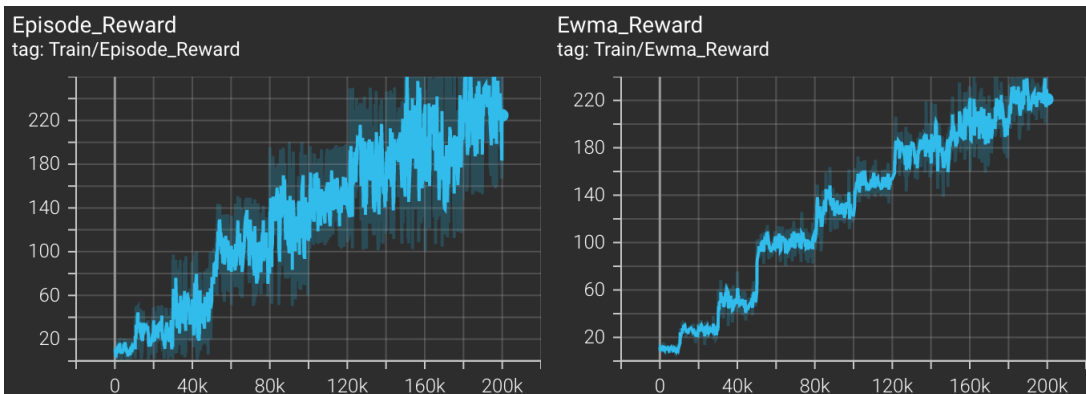


Figure 3.5: Rewards of DQN for *BreakoutNoFrameskip-v4* during training.

```
(dl) root@5b48b3907782:/workspace/lab6# xvfb-run -s "-screen 0 1400x900x24" python dqn_breakout.py --mode test --device 1 --resume_ckpt models/230526_120821_dqn-breakout_batch_size32_episodes20000warmup20000_lr6.25e-05gamma0.99decay1000000min0.1_freq-behavior4freq-target100000freq-eval1000000_eps-eval0.01episodes-eval10/dqn-breakout.pth
Namespace(mode='test', devices=[1], batch_size=32, num_workers=8, capacity=100000, num_train_episodes=20000, num_steps_warmup=20000, lr=6.25e-05, gamma=0.99, epsilon_decay=1000000, epsilon_min=0.1, freq_update_behavior=4, freq_update_target=10000, freq_eval=200000, resume_ckpt=models/230526_120821_dqn-breakout_batch_size32_episodes20000warmup20000_lr6.25e-05gamma0.99decay1000000min0.1_freq-behavior4freq-target100000freq-eval1000000_eps-eval0.01episodes-eval10/dqn-breakout.pth, epsilon_eval=0.01, num_eval_episodes=10, dir_writers='./runs/', dir_models='./models/')

Currently using device cuda:1 ...
Finish loading checkpoint from models/230526_120821_dqn-breakout_batch_size32_episodes20000warmup20000_lr6.25e-05gamma0.99decay1000000min0.1_freq-behavior4freq-target100000freq-eval1000000_eps-eval0.01episodes-eval10/dqn-breakout.pth ...
Start testing ...
Episode 1: 234.48
Episode 2: 239.47
Episode 3: 232.22
Episode 4: 223.81
Episode 5: 246.23
Episode 6: 229.97
Episode 7: 212.54
Episode 8: 207.03
Episode 9: 223.01
Episode 10: 221.91
Average Reward: 227.06
```

Figure 3.6: Rewards of DQN for *BreakoutNoFrameskip-v4* during evaluation.

# Chapter 4

## Discussion

Four questions and their answers are listed in this chapter.

### 4.1 Explain effects of the discount factor

In Equation 4.1,  $\lambda$  is discount factor, and the future effects getting smaller with the degree of  $\lambda$  getting larger.

$$G_t = R_{t+1} + \lambda R_{t+2} + \cdots = \sum_{k=0}^{\infty} \lambda^k R_{t+k+1} \quad (4.1)$$

### 4.2 Explain benefits of epsilon-greedy in comparison to greedy action selection

It's better to balance between explore and exploit with greedy action selection, and sometimes choosing other action to explore may be best action.

### 4.3 Explain the necessity of the target network

With target network and behavior network, the training process can be more stable, since Q target is output from target network and it's updated with lower frequency.

## 4.4 Describe the tricks you used in Breakout and their effects, and how they differ from those used in LunarLander

For *LunarLander-v2*, MSE loss is used, but for *BreakoutNoFrameskip-v4*, smooth L1 loss is used instead (See Listing 4.1).

```
1 def _update_behavior_network(self, gamma):
2     state, action, reward, next_state, done = self._memory.sample(
3         self.batch_size)
4
5     q_value = self._behavior_net(state).gather(dim=1, index=action.long())
6     with torch.no_grad():
7         q_next = self._target_net(next_state).max(dim=1)[0]
8         q_target = reward[:, 0] + gamma * q_next * (1 - done[:, 0])
9
10    loss = F.smooth_l1_loss(q_value, q_target.unsqueeze(1))
11
12    self._optim.zero_grad()
13    loss.backward()
14    for param in self._behavior_net.parameters():
15        param.grad.data.clamp_(-1, 1)
16    self._optim.step()
```

Listing 4.1: Python code of `_update_behavior_network` of DQN for *BreakoutNoFrameskip-v4*.